

# neural

June 19, 2025

```
[ ]: ###

import os
import sys

sys.path.append(os.path.dirname(os.path.dirname(os.path.abspath(__file__))))

import pandas as pd
import torch
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from torch import nn
from torch.utils.data import DataLoader, Dataset

from neural_net.preprocessing import get_train_test_data
# from randomforest import export_model
from utils.graphs import compare, scatter_plot
from utils.neural_utils import LawDataset, NeuralNetwork, predict
```

```
[ ]: # %%

X_train, X_test, y_train, y_test = get_train_test_data("../law_data.csv", ↵
    ↪ "first_pf")

training_data = LawDataset(X_train, y_train)
testing_data = LawDataset(X_test, y_test)

train_dataloader = DataLoader(training_data, batch_size=64, shuffle=True)
test_dataloader = DataLoader(testing_data, batch_size=64, shuffle=False)
```

```
[ ]: # %%

train_features, train_labels = next(iter(train_dataloader))
print(f"Feature batch shape: {train_features.size()}")
print(f"Labels batch shape: {train_labels.size()}")
```

Feature batch shape: torch.Size([64, 25])

Labels batch shape: torch.Size([64])

```
[ ]: # %%  
  
device = torch.accelerator.current_accelerator().type if torch.accelerator.  
    is_available() else "cpu"  
print(f"Using {device} device")
```

Using cuda device

```
[ ]: # %%  
  
model = NeuralNetwork(input_size=X_train.shape[1]).to(device)  
model
```

```
[ ]: NeuralNetwork(  
    (linear_relu_stack): Sequential(  
      (0): Linear(in_features=25, out_features=1, bias=True)  
    )  
)
```

```
[ ]: # %%  
  
learning_rate = 1e-2  
batch_size = 64  
epochs = 10  
  
losses = {}  
  
def train_loop(dataloader, model, loss_fn, optimizer):  
    size = len(dataloader.dataset)  
    model.train()  
    for batch, (X, y) in enumerate(dataloader):  
        X = X.to(device)  
        y = y.to(device).float().view(-1, 1)  
        pred = model(X)  
        # print(X, y)  
        loss = loss_fn(pred, y)  
  
        loss.backward()  
        optimizer.step()  
        optimizer.zero_grad()  
  
        if batch % 100 == 0:  
            loss_value = loss.item()  
            current = batch * batch_size + len(X)  
            print(f"loss: {loss_value:>7f} [{current:>5d}/{size:>5d}]")
```

```

def test_loop(dataloader, model, loss_fn, epoch):
    model.eval()
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    test_loss, correct = 0, 0
    with torch.no_grad():
        for X, y in dataloader:
            X = X.to(device)
            y = y.to(device).float().view(-1, 1)
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            # Pour l'accuracy :
            pred_label = (torch.sigmoid(pred) > 0.5).float()
            correct += (pred_label == y).sum().item()
    test_loss /= num_batches
    correct /= size
    losses[epoch] = test_loss

    scatter_plot(losses, xlabel="Epochs", ylabel="Loss", title="Loss per Epoch")
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss:␣
↪{test_loss:>8f} \n")

```

```

[ ]: # %%

loss_fn = nn.BCEWithLogitsLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

epochs = 10
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train_loop(train_dataloader, model, loss_fn, optimizer)
    test_loop(test_dataloader, model, loss_fn, t)
print("Done!")


```

Epoch 1

```

-----
loss: 0.728525 [ 64/15253]
loss: 0.554184 [ 6464/15253]
loss: 0.451272 [12864/15253]

```



Test Error:


Accuracy: 85.9%, Avg loss: 0.471998

Epoch 2

-----  
loss: 0.479992 [ 64/15253]

loss: 0.413908 [ 6464/15253]

loss: 0.366082 [12864/15253]



Test Error:


Accuracy: 89.1%, Avg loss: 0.370166

Epoch 3

-----  
loss: 0.306903 [ 64/15253]

loss: 0.352540 [ 6464/15253]

loss: 0.273468 [12864/15253]



Test Error:


Accuracy: 89.7%, Avg loss: 0.325475

Epoch 4

-----  
loss: 0.329936 [ 64/15253]

loss: 0.372335 [ 6464/15253]

loss: 0.356452 [12864/15253]




Test Error:

Accuracy: 90.0%, Avg loss: 0.301626

Epoch 5

-----  
loss: 0.308612 [ 64/15253]  
loss: 0.365436 [ 6464/15253]  
loss: 0.306604 [12864/15253]




Test Error:

Accuracy: 90.1%, Avg loss: 0.287642

Epoch 6

-----  
loss: 0.250827 [ 64/15253]  
loss: 0.303448 [ 6464/15253]  
loss: 0.250208 [12864/15253]






Test Error:

Accuracy: 90.1%, Avg loss: 0.278643

Epoch 7

-----  
loss: 0.181801 [ 64/15253]  
loss: 0.258912 [ 6464/15253]  
loss: 0.238733 [12864/15253]




Test Error:

Accuracy: 90.2%, Avg loss: 0.272408

Epoch 8

-----  
loss: 0.247495 [ 64/15253]  
loss: 0.307326 [ 6464/15253]  
loss: 0.289869 [12864/15253]

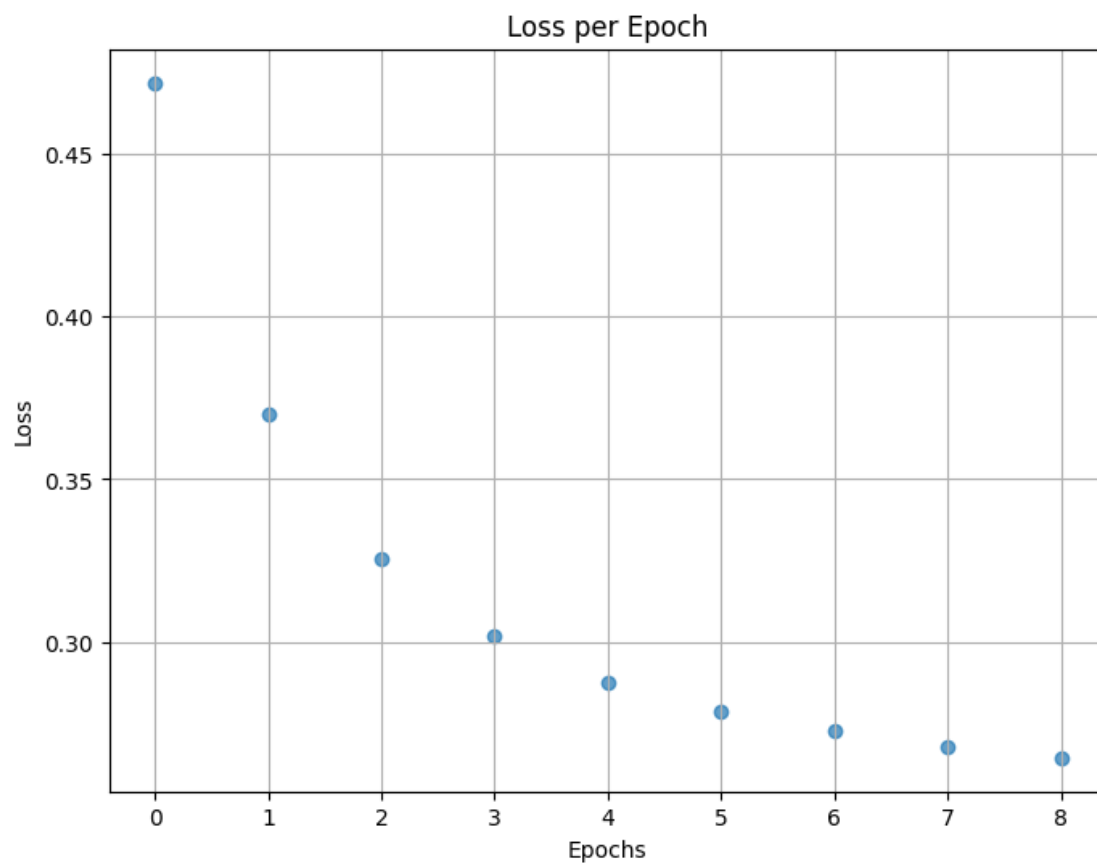


Test Error:

Accuracy: 90.3%, Avg loss: 0.267715

Epoch 9

-----  
loss: 0.295637 [ 64/15253]  
loss: 0.232339 [ 6464/15253]  
loss: 0.266307 [12864/15253]



Test Error:


Accuracy: 90.3%, Avg loss: 0.264333

Epoch 10

-----  
loss: 0.267383 [ 64/15253]

loss: 0.266995 [ 6464/15253]

loss: 0.284033 [12864/15253]



Test Error:

Accuracy: 90.3%, Avg loss: 0.261657

Done!

```
[ ]: # %%  
  
# rf, reg = export_model()  
  
# pred_rf = rf.predict(X_test)  
# pred_reg = reg.predict(X_test)  
  
# # Generate predictions for the neural network on the test set  
# pred_neural = predict(model, X_test)  
  
# compare(  
#     [  
#         classification_report(y_test, pred_rf, output_dict=True),  
#         classification_report(y_test, pred_reg, output_dict=True),
```

```
#         classification_report(y_test, (pred_neural > 0.5).astype(float),  
↪output_dict=True)  
#     ],  
#     model_names=["Random Forest", "Logistic Regression", "Neural Network"],  
#     print_output=True,  
# )
```

```
[ ]: # %%  
  
torch.save(model.state_dict(), "model.pth")
```